

# ANSIBLE T DAILY TASKS

# ANSIBLE INVENTORY MANAGEMENT

- COMPERHENSIVE GUIDE
- THEORY + PRACTICAL TASKS
- REAL TIME SCENARIO TASKS
- END TO END DOC

#### **CONTACT US:**

fenilgajjar.devops@gmail.com linkedin.com/in/fenl-gajjar github.com/Fenil-Gajjar

0000





Unleashed 🧩



#### Welcome to the Ansible Inventory Management Guide

Hello DevOps enthusiasts!

Welcome to this in-depth, practical guide on **Inventory Management in Ansible**— a fundamental yet often underestimated piece of the automation puzzle.

This document isn't just about theory — it's built for **real-world DevOps engineers**, cloud practitioners, and automation fans like you who want to **understand**, **implement**, **and scale** Ansible inventories the right way.

#### What to Expect

In the pages ahead, you'll find:

- Easy-to-understand concepts
- Hands-on practical tasks inspired by real project scenarios
- Clean, YAML-based inventory structures
- Coverage of **static**, **dynamic**, and **mixed-OS** environments
- Industry-backed best practices to build robust, scalable automation

Whether you're just getting started or fine-tuning enterprise-grade inventories, this guide will meet you where you are.

#### **A Note of Thanks**

Before we dive in — a heartfelt thank you to **everyone who has supported me on this journey**.

Your encouragement, feedback, and curiosity have inspired me to build and share this resource. I hope it helps you in your journey just as much as you've helped me in mine.

So, let's dive into the world of Ansible inventory — the place where your infrastructure gets its voice.

Let's automate with clarity, confidence, and best practices.

## **What is Inventory in Ansible?**

In Ansible, **Inventory** refers to the **list of nodes (hosts/servers)** that Ansible manages and automates. It's the **source of truth** that tells Ansible *which machines* to connect to, how to connect to them, and how to group them for efficient management.

Think of it as a dynamic phonebook of your infrastructure.

#### **W** Key Concepts:

- Inventory = Hosts + Groups + Variables
- It defines the target infrastructure: which hosts, their IPs, SSH ports, usernames, etc.
- It can be written in INI, YAML, or dynamically generated via scripts or plugins.

#### **X** Where is Inventory Used?

When running ad-hoc commands:

ansible all -i inventory.ini -m ping

#### When running **playbooks**:

```
ansible-playbook -i inventory.yml site.yml
```

#### Example (INI format):

```
[web]
web01 ansible_host=192.168.1.10 ansible_user=ubuntu
web02 ansible_host=192.168.1.11 ansible_user=ubuntu
[db]
db01 ansible_host=192.168.1.20 ansible_user=postgres
[all:vars]
ansible_ssh_private_key_file=~/.ssh/id_rsa
```

- Groups: [web], [db]
- Host-level variables: ansible\_host, ansible\_user
- Global variables: under [all:vars]

#### Types of Inventories:

- 1. Static Inventory Hardcoded host list (INI or YAML format)
- 2. **Dynamic Inventory** Generated from cloud APIs (AWS, Azure, GCP, etc.)

#### Why is it Important?

- Lets you organize infrastructure by roles, such as web, db, loadbalancer, etc.
- Enables **environment separation** (dev, staging, prod).
- Inventory design impacts how playbooks scale and perform.

## Why Inventory is the Core of Configuration

#### **Management in Ansible**

Inventory is **the backbone** of how Ansible understands and interacts with your infrastructure. Without inventory, **Ansible has no context** — it doesn't know *what* to configure, *where* to connect, or *how* to apply tasks.

Let's break down why it's so fundamental:

#### Defines the Target Infrastructure

Inventory explicitly tells Ansible:

- Which hosts are part of the infrastructure
- How to connect to each host (IP, SSH user, port, key)
- What groups they belong to (e.g., web, db, prod, staging)
- Any host-specific or group-specific variables

Without this information, your playbooks are just recipes with no ingredients.

#### **TENABLES Scalable Configuration**

Inventory allows you to:

- Manage hundreds or thousands of hosts with just a few lines
- Apply configurations based on roles, locations, environments
- Control rolling updates, conditional logic, and host selection

This scalability is essential in large-scale infrastructure.

#### Dynamic Adaptability

With dynamic inventory:

- Your infrastructure can scale up/down automatically
- Ansible queries cloud APIs (e.g., AWS EC2, Azure VM) in real time
- You don't need to manually update inventory files

This is critical in cloud-native and containerized environments.

#### 🧩 Drives Variable Management

Inventory isn't just a host list — it's a place to **store and organize variables**, like:

- Hostnames, IPs, connection details
- Application ports, database credentials
- Environment-specific overrides

This tight integration with **group\_vars** and **host\_vars** makes configuration clean, reusable, and secure.

#### ☐ Core for Targeting and Execution

Every Ansible command — whether ad-hoc or via playbook — begins with a question:

"Which hosts should I run this against?"

Inventory answers that.

**Examples:** 

```
ansible web -i inventory.ini -m ping
ansible-playbook -i inventory.yml deploy.yml --limit db
```

Without the inventory, these commands are meaningless.

#### X

## Types of Inventories in Ansible: Static vs Dynamic

In Ansible, inventories come in two main forms — **Static** and **Dynamic**. Both serve the same purpose (defining the hosts Ansible will manage), but they're suited for different kinds of environments.

Let's explore each type with explanations, examples, and real-world use cases.

#### **Static Inventory**



A manually defined list of hosts and groups, written in formats like INI or YAML.

#### Common file locations:

- inventory.ini
- inventory.yml
- Passed directly using -i flag in commands

```
Example: INI Format
[web]
web01 ansible_host=192.168.1.10
web02 ansible_host=192.168.1.11
[db]
db01 ansible_host=192.168.1.20
[all:vars]
ansible_user=ubuntu
ansible_ssh_private_key_file=~/.ssh/id_rsa
Example: YAML Format
all:
  children:
    web:
      hosts:
        web01:
```

ansible\_host: 192.168.1.10

```
web02:
    ansible_host: 192.168.1.11

db:
    hosts:
    db01:
        ansible_host: 192.168.1.20

vars:
    ansible_user: ubuntu
    ansible_ssh_private_key_file: ~/.ssh/id_rsa
```

#### **Best for:**

- Small to medium-sized environments
- Static infrastructure (bare metal, small VMs)
- Dev/testing environments

#### X Limitations:

• Needs manual updates when hosts change

• Doesn't scale well in cloud or dynamic environments



★ What is it?

An **auto-generated list of hosts**, pulled in real-time from **external sources** like cloud providers, CMDBs, or APIs.

Instead of hardcoding host entries, Ansible uses **inventory plugins** or scripts to **query APIs** and get the latest data.

#### **How it works:**

- You define a **plugin config file** (usually YAML)
- Ansible pulls live data from sources like AWS, GCP, Azure, VMware, etc.
- Supports **filters**, **host variables**, **tags**, etc.

Example: AWS EC2 Dynamic Inventory

plugin: aws\_ec2

regions:

- us-east-1

filters:

```
tag:Environment: production
```

#### keyed\_groups:

- key: tags.Name

#### hostnames:

- public-ip-address

This fetches all EC2 instances in us-east-1 with the Environment=production tag.

#### **Best for:**

- Cloud-native environments
- Auto-scaling groups
- Infrastructure with frequent changes
- Multi-region or multi-cloud architectures

#### Supported sources:

- AWS (aws\_ec2)
- Azure (azure\_rm)

- GCP (gcp\_compute)
- Kubernetes (k8s)
- VMware (vmware\_vm\_inventory)
- Custom scripts (using script plugin)

#### X Limitations:

- Requires initial setup and authentication configs
- May introduce **latency** when fetching large inventories
- Requires Ansible >= 2.8 for full plugin support

## Inventory File Basics in Ansible

(How to write and structure a static inventory file)

In Ansible, the inventory file is the **starting point** of every automation — it tells Ansible which hosts to manage and how to group and configure them. A static inventory is typically written in INI or YAML format, with YAML being the more modern and flexible choice.

#### Where Is the Inventory File Located?

You can:

Pass it explicitly using the -i option:

ansible-playbook -i inventory.yml site.yml

Set a **default inventory file** in ansible.cfg:

[defaults]

inventory = ./inventory/inventory.yml

Typical file locations:

## **INI Format (Legacy, Simple)**

A traditional INI inventory looks like this:

```
[web]
web01 ansible_host=192.168.1.10 ansible_user=ubuntu
web02 ansible_host=192.168.1.11 ansible_user=ubuntu
[db]
db01 ansible_host=192.168.1.20 ansible_user=postgres
```

```
[all:vars]
ansible_ssh_private_key_file=~/.ssh/id_rsa
```

Pros: Simple, quick to write

X Cons: Not very expressive or scalable for complex setups

#### **YAML Format (Modern, Recommended)**

YAML allows for **structured**, **nested**, and **more readable** inventory files. It's supported since Ansible 2.4+ and is perfect for advanced use cases.

#### Basic Example:

```
all:
  hosts:
  web01:
    ansible_host: 192.168.1.10
    ansible_user: ubuntu
  web02:
    ansible_host: 192.168.1.11
    ansible_user: ubuntu
  children:
```

```
db:
       hosts:
         db01:
           ansible_host: 192.168.1.20
           ansible_user: postgres
  • all: root group, required
  • hosts: individual hosts
  • children: nested groups like web, db, etc.
With Group Variables:
all:
  children:
    web:
       hosts:
         web01:
         web02:
       vars:
```

```
app_port: 8080
    db:
      hosts:
        db01:
      vars:
        ansible_user: postgres
        db_port: 5432
With Host Variables:
all:
  hosts:
    web01:
      ansible_host: 192.168.1.10
      ansible_port: 2222
      ansible_user: deployer
```

ansible\_user: ubuntu

#### 4. Directory Structure Best Practices

For scalable environments, organize your inventory like this:

This lets you **separate config from structure**, making it easier to manage complex setups.

#### Test and Visualize Your Inventory

Use the ansible-inventory command:

# Outputs the entire inventory in JSON format (expanded view)

ansible-inventory -i inventory.yml --list

# Displays a visual tree-like structure of groups and hosts ansible-inventory -i inventory.yml --graph

## X Task: Define a Basic YAML Inventory for a

#### **Master-Worker Infrastructure on AWS**

#### Scenario Overview:

You're managing a small cloud-based Kubernetes-like infrastructure on **AWS**, consisting of:

- 1 Master Node responsible for orchestration and control
- 3 Worker Nodes responsible for running workloads (containers, apps, etc.)

This setup needs to be clearly defined in your **Ansible inventory file**, with proper grouping and variables to support automated provisioning or configuration.

#### **Objective:**

Create a clean, readable, and well-structured **YAML-based static inventory** for this infrastructure. The file should:

- Group all the worker nodes under workers
- Include the master in its own group

- Provide relevant host-specific details (like ansible\_host, ansible\_user)
- Define common variables at the group level

#### Directory Assumption:

We'll assume this inventory is in the following path:

```
inventory/
└─ inventory.yml
inventory.yml (YAML Static Inventory File)
all:
  children:
    master:
      hosts:
        master-node:
           ansible_host: 3.91.102.5
```

ansible\_user: ec2-user

```
ansible_ssh_private_key_file: ~/.ssh/aws_key.pem
workers:
  hosts:
    worker-node-1:
      ansible_host: 3.88.210.23
      ansible_user: ec2-user
      ansible_ssh_private_key_file: ~/.ssh/aws_key.pem
    worker-node-2:
      ansible_host: 3.92.17.44
      ansible_user: ec2-user
      ansible_ssh_private_key_file: ~/.ssh/aws_key.pem
    worker-node-3:
      ansible_host: 3.93.84.112
      ansible_user: ec2-user
```

ansible\_ssh\_private\_key\_file: ~/.ssh/aws\_key.pem

#### vars:

```
ansible_connection: ssh
```

ansible\_port: 22

#### Explanation of Key Sections

- **Groups (master, workers)**: Organize hosts by their function for targeted automation.
- ansible\_host: Public IP of each AWS EC2 instance.
- ansible\_user: Default EC2 user (e.g., ec2-user for Amazon Linux).
- ansible\_ssh\_private\_key\_file: Path to your SSH key to connect.
- Global variables (all.vars): Apply to every host SSH connection, default port.

#### **★** Bonus Tips:

You can now run a ping test on just worker nodes:

ansible workers -i inventory/inventory.yml -m ping

• Or run a full cluster health check using a playbook:

ansible-playbook -i inventory/inventory.yml
playbooks/health-check.yml

#### What Are Inventory Groups and Host Variables?

#### In Ansible:

- **Groups** are *collections of hosts* with similar roles (e.g. web, db, prod, workers).
- **Host variables** are settings that apply to a specific host.
- **Group variables** are settings that apply to every host in that group.
- You can even **nest groups** using children.

Inventory is more than just a list — it's a blueprint of your infrastructure, where grouping and variables help you manage with precision and elegance.

#### Why Do They Matter?

Using groups and variables:

- Makes playbooks modular and reusable
- Avoids **repetition** of SSH settings, credentials, or role-based parameters
- Supports role-specific logic, host-level overrides, and environmental separation

#### Anatomy of a Well-Structured Inventory (YAML Style)

Here's a rich, real-world example of a YAML inventory that demonstrates:

- Group-level variables
- Host-level variables
- Group nesting (via children)
- Multiple environments
- Clean separation of logic

```
Example: Full YAML Inventory Structure
```

```
all:
    children:
    production:
        children:
        web:
        hosts:
        web1.prod.example.com:
```

```
ansible_host: 34.201.100.10
      ansible_user: ubuntu
    web2.prod.example.com:
      ansible_host: 34.201.100.11
      ansible_user: ubuntu
 vars:
    app_env: production
    nginx_port: 80
db:
 hosts:
    db1.prod.example.com:
      ansible_host: 10.0.1.10
      ansible_user: postgres
      db_engine: postgresql
 vars:
    db_port: 5432
    db_user: admin
```

```
staging:
    children:
      web:
        hosts:
          web1.staging.example.com:
            ansible_host: 34.203.100.20
            ansible_user: ubuntu
        vars:
          app_env: staging
          nginx_port: 8080
vars:
  ansible_ssh_private_key_file: ~/.ssh/aws-key.pem
  ansible_connection: ssh
  ansible_port: 22
```

## Let's Break It Down:

#### 1. all Group

This is the **top-most group**, *automatically defined* by Ansible. It includes everything.

yaml

CopyEdit

all:

#### 2. Nested Environments: production and staging

These are **environmental groupings**, commonly used in cloud setups:

production:

children:

web:

db:

- production includes web and db
- Same structure for staging

This lets you run playbooks like:

```
ansible-playbook -i inventory.yml site.yml --limit
production
```

#### 3. Host-Specific Variables

These override group or global settings and are used for per-host tweaks:

```
db1.prod.example.com:
   ansible_host: 10.0.1.10
   db_engine: postgresql
```

#### **4.** Group Variables (vars)

Define shared settings for a group (no need to repeat on each host!):

```
vars:
   app_env: production
   nginx_port: 80
```

All hosts in the web group will inherit this.

#### 5. Global Variables

Defined under all.vars, this applies to **every host** in the inventory:

```
all:
```

```
vars:
```

```
ansible_connection: ssh
ansible_ssh_private_key_file: ~/.ssh/aws-key.pem
```

#### Great for:

- SSH key config
- Port setup
- Connection settings

#### BONUS: External Variables (Cleaner Design)

When your inventory grows, you can move variables out of the inventory file for clarity:

Now you can clean up your inventory.yml and let Ansible auto-load these variable files.

#### **Tips for Best Practices**

Tip	Why it matters
Use groups to reduce repetition	One setting for many hosts
Use group_vars/ and host_vars/ folders	Cleaner, scalable, modular
Never hardcode secrets in inventory	Use Ansible Vault instead
Override wisely	Host-level vars override group-level vars
See Document groups clearly	Helps collaboration in team environments

#### Why Use Host Ranges?

Imagine you have 10 worker nodes named like this:

```
worker01, worker02, ..., worker10
```

Instead of writing each one manually, Ansible lets you define them as a range, saving time and keeping your inventory DRY (Don't Repeat Yourself).

### INI Format (Legacy, but simple)

```
[workers]
worker[01:10].example.com ansible_user=ubuntu
worker[01:10].example.com will expand to:
worker01.example.com
worker02.example.com
worker10.example.com
```

ightharpoonup Tip: Leading zero (01) ensures numbers are zero-padded.

#### YAML Format (Recommended)

In YAML, you **can't** directly use the worker[01:10] syntax — but you can still keep it readable using a Jinja loop with an **external script or dynamic inventory**, or just define it clearly with list-style formatting.

But for static inventories in YAML, define explicitly like this:

```
all:
    children:
    workers:
    hosts:
        worker01.example.com:
        ansible_user: ubuntu
        worker02.example.com:
        ansible_user: ubuntu
        worker03.example.com:
        ansible_user: ubuntu
        worker04.example.com:
```

```
ansible_user: ubuntu
worker05.example.com:
ansible_user: ubuntu
```

Pro Tip: You can generate the YAML snippet programmatically using Python or Bash for dozens/hundreds of hosts.

#### X Automation Trick: Generate with Python

```
for i in range(1, 11):
    print(f" worker{str(i).zfill(2)}.example.com:")
    print(" ansible_user: ubuntu")
```

Paste the result directly into your inventory.yml.

#### Assigning Variables to Hosts and Groups in Ansible

When you're managing infrastructure at scale, you often need to assign variables like:

- SSH user
- Port numbers
- Environment names
- Custom app config
- Credentials (preferably encrypted with Vault)

In Ansible, you assign variables at two key levels:

#### **1.** Group Variables → Apply to Many Machines

Use **group variables** when you want **all hosts in a group** (e.g. all web servers or all db nodes) to share the same config.

When to use: All web servers use nginx, or all staging nodes share the same environment variable.

#### **Example: Assigning Group Variables (Inside Inventory YAML)**

```
all:
    children:
    web:
    hosts:
        web1.example.com:
        web2.example.com:
    vars:
        ansible_user: ubuntu
        app_env: production
        nginx_port: 80
```

All hosts in the web group:

- Use the ubuntu user
- Serve the production environment
- Run Nginx on port 80

#### Preferred (Clean) Method: Use group\_vars/

```
inventory/

inventory.yml

group_vars/
web.yml
```

#### group\_vars/web.yml:

ansible\_user: ubuntu

app\_env: production

nginx\_port: 80

Ansible automatically loads this when the web group is targeted.

#### @ 2. Host Variables → Apply to One Machine

Use host variables when a host has custom settings that shouldn't affect others.

When to use: One server has a different IP, custom SSH port, or database credentials.

**Example: Assigning Host Variables** 

```
all:
   hosts:
   db1.example.com:
     ansible_host: 192.168.10.10
     ansible_user: postgres
```

db\_engine: postgresql

db\_port: 5432

Only  $db1.example.com\ will\ receive\ these\ values.$ 

#### Preferred Method: Use host\_vars/

inventory/

inventory.yml

host\_vars/

db1.example.com.yml

#### host\_vars/db1.example.com.yml:

ansible\_user: postgres

db\_engine: postgresql

db\_port: 5432

## Variable Precedence

Level	Priority	
Host variables	▲ Highest	
Group variables	Lower	
all group vars	Even lower	
Defaults in	Lowest	

## **Connecting to Hosts: Behavioral Inventory Parameters in Ansible**

#### **®** What Are Behavioral Inventory Parameters?

Behavioral parameters are **Ansible-specific variables** that control *how* it connects to your hosts — over SSH, WinRM, custom ports, different usernames, private keys, etc.

You assign these variables either:

- In the inventory file
- In host\_vars/ or group\_vars/
- Or pass them via CLI or playbooks

#### Commonly Used Behavioral Parameters

Here's a **comprehensive cheat sheet** of useful behavioral parameters — all with real-world examples.

- ansible\_host
- Defines the real IP or DNS of the host (if hostname is just a label)

```
web1:
   ansible_host: 192.168.1.10
```

Useful when your internal DNS name differs from the SSH target.

- ansible\_user
- Username to connect with (e.g., ubuntu, ec2-user, root)

```
ansible_user: ubuntu
```

Overrides your default local user. Needed for cloud machines.

- ansible\_port
- SSH Port (default is 22)

```
ansible_port: 2222
```

If your SSH service runs on a custom port.

- ansible\_ssh\_private\_key\_file
- Path to private SSH key

```
ansible_ssh_private_key_file: ~/.ssh/aws-key.pem
```

Crucial for connecting to cloud hosts like AWS EC2.

- ansible\_connection
- Connection type (usually ssh, but also local, docker, winrm, etc.)

```
ansible_connection: ssh
```

Use local for localhost or docker when managing containers.

- ansible\_become, ansible\_become\_user
- Enable privilege escalation (like sudo)

```
ansible_become: true
ansible_become_user: root
```

Needed when you connect as a non-root user but need root privileges.

- ansible\_python\_interpreter
- Set specific Python path (important for non-default distros or environments)

```
ansible_python_interpreter: /usr/bin/python3
```

Fixes issues when /usr/bin/python isn't available or isn't Python 3.

#### Example YAML Inventory with Behavioral Parameters

```
all:
   children:
   web:
```

hosts:

# web1.example.com: ansible\_host: 10.0.0.11 ansible\_user: ubuntu ansible\_port: 22 ansible\_ssh\_private\_key\_file: ~/.ssh/aws-key.pem ansible\_connection: ssh ansible\_become: true ansible\_become\_user: root ansible\_python\_interpreter: /usr/bin/python3

This setup ensures Ansible connects, escalates privileges, and uses the correct Python on web1.

## Testing Host Connectivity

Use the ping module to validate everything's working:

```
ansible all -i inventory.yml -m ping
```

Or test a single host:

```
ansible web1.example.com -i inventory.yml -m ping -u ubuntu
```

### Best Practices for Inventory Management in Ansible

These best practices apply whether you're managing 10 nodes or 10,000 — across on-prem, cloud, hybrid, or containerized environments.

#### 1 Use YAML (INI is legacy)

**Why:** YAML is structured, easier to read, and supports nested groups and variables cleanly.

Recommended:

all:

```
children:
   web:
   hosts:
    web1.example.com:
```

web2.example.com:

```
X Avoid:
```

```
[web]
web1.example.com
```

```
web2.example.com
```

YAML promotes consistency and integrates better with modern tools.

#### 2 Group by Function, Role, or Environment

Use logical groupings like:

```
• web, db, cache
```

```
• dev, staging, prod
```

• frontend, backend

```
Example:
```

all:

```
children:
```

prod:

children:

web:

db:

Makes targeting playbooks easy and readable: ansible-playbook -1 web site.yml

#### 3 Use group\_vars/ and host\_vars/ Instead of Inline Variables

Keep your inventory clean and separate config details.

Recommended:

Improves clarity, reuse, and security (with Vault).

#### 4 Use Descriptive Hostnames (Not IPs Directly)

Use DNS names or aliases in inventory — assign IPs via ansible\_host.

**Better:** 

#### web1:

```
ansible_host: 192.168.1.10
```

X Avoid:

192.168.1.10

You gain flexibility if the IP changes.

#### **5** Use Dynamic Inventory for Cloud Infrastructure

Static inventories are not scalable in cloud-native environments.

**U**se:

- AWS: aws\_ec2 plugin
- GCP: gcp\_compute

• Azure: azure\_rm

```
ansible-inventory -i aws_ec2.yml --list
```

Dynamic inventories sync automatically with your live infra.

#### 6 Avoid Duplicating Hosts in Multiple Groups

Avoid placing the same host in multiple unrelated groups — it causes variable conflicts and logic bugs.

#### **Better:**

- Use nested groups if needed
- Or clearly isolate roles

#### Always Set the Python Interpreter (Especially for Minimal OS)

Set:

```
ansible_python_interpreter: /usr/bin/python3
```

Prevents python not found errors on distros like Ubuntu Minimal, Amazon Linux 2, etc.

#### 8 Use Vault for Sensitive Variables

Never store plain passwords, keys, or secrets in the inventory.

**U**se:

```
ansible-vault encrypt_string --name 'db_password'
'SuperSecret123'
```

Or encrypt entire group\_vars/prod.yml.

#### 9 Use Behavioral Variables Only Where Needed

Avoid cluttering inventory with repeated values like ansible\_user, ansible\_port, etc., if the defaults work.

✓ Use at group level (e.g. for all web servers) instead of per-host.

#### 10 Document Inventory Structure Clearly

Maintain a README inside your inventory directory explaining:

- Group naming conventions
- Host naming logic

• Inventory split strategy (per env, team, app)

Makes collaboration and onboarding smoother.

## **\*** Bonus Tips:

Practice	Why It Matters
Keep inventory in Git	Version control & rollback
Tag servers consistently	Useful with dynamic inventory plugins
Use inventory plugins over scripts	Native, safer, and better maintained
Test inventory syntax	Use ansible-inventorygraph orlist

# **X** Task: Configure Host-Specific SSH Ports in Ansible Inventory

#### Scenario

You're managing a fleet of Linux servers, and for security reasons, each host uses a **custom SSH port** instead of the default 22.

Here's the infrastructure:

Hostname	IP Address	SSH Port
web1.example .com	192.168.1 0.10	2222
db1.example.	192.168.1 0.20	2200
cache1.examp	192.168.1 0.30	2022

You want Ansible to connect to each host using its correct SSH port.

#### Step 1: Create YAML Inventory File

```
all:
  children:
    web:
      hosts:
        web1.example.com:
          ansible_host: 192.168.10.10
          ansible_port: 2222
          ansible_user: ubuntu
          ansible_ssh_private_key_file: ~/.ssh/id_rsa_web
    db:
      hosts:
        db1.example.com:
          ansible_host: 192.168.10.20
          ansible_port: 2200
          ansible_user: ubuntu
          ansible_ssh_private_key_file: ~/.ssh/id_rsa_db
    cache:
      hosts:
```

#### cache1.example.com:

ansible\_host: 192.168.10.30

ansible\_port: 2022

ansible\_user: ubuntu

ansible\_ssh\_private\_key\_file: ~/.ssh/id\_rsa\_cache

#### **Explanation**

- ansible\_host: The actual IP of the machine
- ansible\_port: The custom SSH port Ansible should use to connect
- ansible\_user: Login user on the remote machine
- ansible\_ssh\_private\_key\_file: Key specific to that host

⚠ If you're using the same SSH key for all hosts, you can move ansible\_ssh\_private\_key\_file to group\_vars/all.yml instead.

#### Step 2: Test Connections

Ping all hosts to ensure Ansible can reach them via the correct port:

```
ansible all -i inventory.yml -m ping
```

Or test a specific group:

```
ansible web -i inventory.yml -m ping
```

#### **®** What This Teaches

- How to use host-specific SSH ports with ansible\_port
- Clean separation of infrastructure by function (web, db, cache)
- How to structure an inventory with real host metadata

# **X** Task: Define Multiple Environments (dev, staging, prod) in a Single Inventory

#### Scenario

Your team manages three environments:

- dev: For internal development
- staging: For QA testing
- prod: For live users

Each environment has its own set of app servers and DBs, but **uses the same playbook**. You want to:

- Structure your inventory cleanly by environment
- Make it easy to target specific environments
- Assign shared variables to each environment (like region, user, interpreter)

#### Step 1: Inventory Layout (inventory.yml)

```
all:
  children:
    dev:
      children:
        dev_app:
          hosts:
            dev-app1.example.com:
              ansible_host: 10.0.1.10
            dev-app2.example.com:
              ansible_host: 10.0.1.11
        dev_db:
          hosts:
            dev-db1.example.com:
              ansible_host: 10.0.1.20
    staging:
      children:
```

```
staging_app:
      hosts:
        staging-app1.example.com:
          ansible_host: 10.0.2.10
    staging_db:
      hosts:
        staging-db1.example.com:
          ansible_host: 10.0.2.20
prod:
  children:
    prod_app:
      hosts:
        prod-app1.example.com:
          ansible_host: 10.0.3.10
        prod-app2.example.com:
          ansible_host: 10.0.3.11
    prod_db:
      hosts:
```

#### prod-db1.example.com:

ansible\_host: 10.0.3.20

#### Step 2: Add group\_vars/ for Each Environment

Create files like:

group\_vars/dev.yml

ansible\_user: ubuntu

ansible\_python\_interpreter: /usr/bin/python3

region: us-west-1

env: dev

#### group\_vars/prod.yml

ansible\_user: ansibleadmin

ansible\_python\_interpreter: /usr/bin/python3

region: us-east-1

env: prod

This way, each environment gets its own SSH user, region config, etc., without cluttering the main inventory.

#### Step 3: Target Environments in Playbooks

To run your playbook only on dev:

```
ansible-playbook -i inventory.yml site.yml -l dev
```

To run only on production DBs:

```
ansible-playbook -i inventory.yml site.yml -l prod_db
```

#### **®** What You Learn

- Structuring inventory around environments
- Using nested groups for apps vs databases
- Keeping environment-specific config in group\_vars/
- Dynamically targeting any layer: full env, app layer, or DB layer

# **X** Task: Manage a Mixed Environment of Linux and Windows Hosts in a Single Inventory

#### Scenario

You're managing an environment that includes:

- Linux servers for application hosting (Ubuntu or CentOS)
- Windows servers for Active Directory and file sharing

#### You want to:

- Manage both OS types with the same Ansible inventory
- Use the right connection methods (ssh vs winrm)
- Assign correct interpreters, users, ports, and privilege settings

#### Step 1: Define Inventory (inventory.yml)

```
all:
  children:
    linux_servers:
      hosts:
        app01.linux.local:
          ansible_host: 192.168.10.10
        db01.linux.local:
          ansible_host: 192.168.10.11
   windows_servers:
      hosts:
        win01.windows.local:
          ansible_host: 192.168.10.20
        ad01.windows.local:
          ansible_host: 192.168.10.21
```

#### Step 2: Set Group Variables

#### group\_vars/linux\_servers.yml

```
ansible_connection: ssh
ansible_user: ubuntu
ansible_become: true
ansible_become_method: sudo
ansible_python_interpreter: /usr/bin/python3
```

#### group\_vars/windows\_servers.yml

```
ansible_connection: winrm
ansible_user: Administrator
ansible_password: "{{ vault_windows_password }}"
ansible_port: 5986
ansible_winrm_transport: basic
ansible_winrm_server_cert_validation: ignore
```

**Tip**: Use Ansible Vault for passwords like vault\_windows\_password.

#### Representation of the strain o

Use PowerShell on the Windows host to enable WinRM (simplified):

```
winrm quickconfig
winrm set winrm/config/service/Auth @{Basic="true"}
winrm set winrm/config/service @{AllowUnencrypted="true"}
Enable-PSRemoting -Force
```

Or better, use an Ansible role like ansible-windows-winrm if bootstrapping remotely.

#### Step 4: Test Connectivity

ansible all -i inventory.yml -m ping

- You should see pong from both Linux and Windows hosts
- If any fail, try -vvv for verbose output

#### S Bonus: Mixed OS Playbook Targeting

You can now write OS-specific tasks and use when clauses:

- name: Mixed OS playbook
hosts: all
tasks:
 - name: Linux-specific task
 shell: uname -a
 when: ansible\_connection == 'ssh'

 - name: Windows-specific task
 win\_command: hostname
 when: ansible\_connection == 'winrm'

#### **@** What You Learn

- How to manage mixed environments cleanly
- Differentiate connection types (ssh, winrm)
- Set appropriate interpreters, ports, users
- Structure inventory for OS-specific logic
- Use Ansible Vault to protect sensitive data

#### Wrap-Up: From Inventory Chaos to Inventory Clarity

And that's a wrap! 🎉

You've just completed a deep, real-world journey through one of the most crucial pillars of Ansible — **Inventory Management**.

In this doc, we explored how to:

- Understand what an inventory is and why it's the heart of Ansible
- Structure clean, scalable inventories using YAML
- Use static and dynamic inventories effectively
- Apply real-time scenarios like multi-environment setups and mixed OS infrastructure
- Implement behavioral parameters, best practices, and secure connections
- Bring clarity, order, and professionalism to your infrastructure automation

This wasn't just a theory dump — it was crafted with real project scenarios to help **you implement, not just learn**.

#### Thanks for Being Part of This

I truly appreciate your support and time in reading through this guide.

Your encouragement and interest keep this journey going — and I hope this document added real value to your DevOps toolbox.

If this helped you or sparked any new ideas, feel free to share it, drop feedback, or connect further.

#### What's Next?

Stay tuned!

The next doc in this series will dive into:

#### 

- How to assign variables cleanly and flexibly
- The difference between group vars, host vars, and defaults
- Real use-cases and directory structure tips

So if you enjoyed this guide — the next one is going to level you up even more.



~ ' Fenil Gajjar ':

Let's keep building clean, efficient, and intelligent automation — one YAML file at a time. 🙌 "

<sup>&</sup>quot; Thanks again for walking this path with me.